

```
In [1]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.arima_model import ARIMA
from pmdarima.arima import auto_arima
from sklearn.metrics import mean_squared_error, mean_absolute_error
import math
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neural_network import MLPRegressor
from sklearn.metrics import mean_squared_error, r2_score
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.callbacks import EarlyStopping
```

```
In [2]: df = pd.read_csv('dataset/train.csv')
submit = pd.read_csv('dataset/sample_submission.csv')
X = pd.read_csv('dataset/Weather.csv')
solar = pd.read_excel('dataset/solar_irradiance_full.xlsx')
# Convert the 'Timestamp' to a datetime object
df['Timestamp'] = pd.to_datetime(df['Timestamp'], format='%b %d, %Y %I%p')

# Extract the minimum and maximum dates from the original data
min_date = df['Timestamp'].min().date()
max_date = df['Timestamp'].max().date()

# Create a dictionary from the original data to map timestamps to baseline values
timestamp_baseline_map = dict(zip(df['Timestamp'], df['% Baseline']))

# Create a list to hold the new DataFrame rows
new_rows = []

# Iterate over each date from min_date to max_date
for date in pd.date_range(start=min_date, end=max_date, freq='D'):
    # Generate hourly timestamps for the specific date starting from 7 AM
    hourly_timestamps = pd.date_range(start=f'{date} 07:00:00', end=f'{date} 23:00:00')

    # Add the new rows with values from the map or a placeholder if missing
    for timestamp in hourly_timestamps:
        baseline_value = timestamp_baseline_map.get(timestamp, np.nan) # Use 0 if
        new_rows.append([timestamp, baseline_value])

# Create the final DataFrame
final_df = pd.DataFrame(new_rows, columns=['Timestamp', '% Baseline'])

# Find the last timestamp in the original data
last_timestamp = df['Timestamp'].max()

# Filter the final_df to only include timestamps up to the last_timestamp in the or
df_filter = final_df[final_df['Timestamp'] <= last_timestamp]
```

```

# Convert both columns to datetime objects
X['date_time'] = pd.to_datetime(X['date_time'], format='%m/%d/%Y %H:%M')
df_filter.loc[:, 'Timestamp'] = pd.to_datetime(df_filter['Timestamp'])

# Format both columns to have the same string representation
X.loc[:, 'date_time'] = X['date_time'].dt.strftime('%Y-%m-%d %H:%M:%S')
df_filter.loc[:, 'Timestamp'] = df_filter['Timestamp'].dt.strftime('%Y-%m-%d %H:%M:

# Combine the Year, Month, Day, Hour, and Minute columns to create a Timestamp column
solar['Timestamp'] = pd.to_datetime(solar[['Year', 'Month', 'Day', 'Hour', 'Minute']]
solar.drop(['Year', 'Month', 'Day', 'Hour', 'Minute'], axis=1, inplace=True)

```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\1225792419.py:39: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_filter.loc[:, 'Timestamp'] = pd.to_datetime(df_filter['Timestamp'])
```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\1225792419.py:39: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values in place instead of always setting a new array. To retain the old behavior, use either `df[df.columns[i]] = newvals` or, if columns are non-unique, `df.isetitem(i, newvals)`

```
df_filter.loc[:, 'Timestamp'] = pd.to_datetime(df_filter['Timestamp'])
```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\1225792419.py:42: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values in place instead of always setting a new array. To retain the old behavior, use either `df[df.columns[i]] = newvals` or, if columns are non-unique, `df.isetitem(i, newvals)`

```
X.loc[:, 'date_time'] = X['date_time'].dt.strftime('%Y-%m-%d %H:%M:%S')
```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\1225792419.py:43: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_filter.loc[:, 'Timestamp'] = df_filter['Timestamp'].dt.strftime('%Y-%m-%d %H:%M:%S')
```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\1225792419.py:43: DeprecationWarning: In a future version, `df.iloc[:, i] = newvals` will attempt to set the values in place instead of always setting a new array. To retain the old behavior, use either `df[df.columns[i]] = newvals` or, if columns are non-unique, `df.isetitem(i, newvals)`

```
df_filter.loc[:, 'Timestamp'] = df_filter['Timestamp'].dt.strftime('%Y-%m-%d %H:%M:%S')
```

In [3]: `solar.head()`

Out[3]:

	DHI	DNI	GHI	Clearsky DHI	Clearsky DNI	Clearsky GHI	Cloud Type	Dew Point	Solar Zenith Angle	Surface Albedo	Wind Speed
0	0.0	0.0	0.0	0.0	0.0	0.0	Probably Clear	-6	124.02	0.12	3.5
1	0.0	0.0	0.0	0.0	0.0	0.0	Probably Clear	-5	135.09	0.12	3.8
2	0.0	0.0	0.0	0.0	0.0	0.0	Probably Clear	-5	145.77	0.12	4.2
3	0.0	0.0	0.0	0.0	0.0	0.0	Probably Clear	-6	155.07	0.12	4.5
4	0.0	0.0	0.0	0.0	0.0	0.0	Probably Clear	-8	160.55	0.12	4.6

In [4]: `df_filter.head()`

Out[4]:

	Timestamp	% Baseline
0	2014-01-01 07:00:00	0.0079
1	2014-01-01 08:00:00	0.1019
2	2014-01-01 09:00:00	0.3932
3	2014-01-01 10:00:00	0.5447
4	2014-01-01 11:00:00	0.5485

```
In [5]: # Filter X based on matching timestamps in df_filter
X_filtered = X[X['date_time'].isin(df_filter['Timestamp'])]
X_filtered.drop(['moonrise', 'moonset', 'sunrise', 'sunset'], axis=1, inplace=True)

# Filter Solar
solar_filtered = solar[solar['Timestamp'].isin(df_filter['Timestamp'])]
```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\1412383839.py:3: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
X_filtered.drop(['moonrise', 'moonset', 'sunrise', 'sunset'], axis=1, inplace=True)
```

```
In [6]: # Perform the merge/join based on the timestamp columns
join = pd.merge(df_filter, X_filtered, left_on='Timestamp', right_on='date_time')
# Ensure both 'Timestamp' columns are in datetime format
join['Timestamp'] = pd.to_datetime(join['Timestamp'])
solar_filtered['Timestamp'] = pd.to_datetime(solar_filtered['Timestamp'])

# Perform the merge on the 'Timestamp' column
```

```

join = pd.merge(join, solar_filtered, on='Timestamp', how='inner')
# Drop the redundant 'date_time' column after merge if desired
join = join.drop(columns=['date_time', 'Timestamp'])
join_fil = join.dropna()
join_fil.head()

```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\2765218598.py:5: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```


solar_filtered['Timestamp'] = pd.to_datetime(solar_filtered['Timestamp'])

```

Out[6]:

	% Baseline	maxtempC	mintempC	totalSnow_cm	sunHour	uvIndex	moon_illumination
0	0.0079	-3	-6	0.0	8.7	2	1
1	0.1019	-3	-6	0.0	8.7	2	1
2	0.3932	-3	-6	0.0	8.7	2	1
3	0.5447	-3	-6	0.0	8.7	2	1
4	0.5485	-3	-6	0.0	8.7	2	1

5 rows × 8 columns



In [7]:

```

# Calculate the correlation matrix
correlation_matrix = join_fil.corr()

# Create a heatmap using seaborn
plt.figure(figsize=(22, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f')

# Set the title and display the plot
plt.title('Correlation Matrix Heatmap')
plt.show()

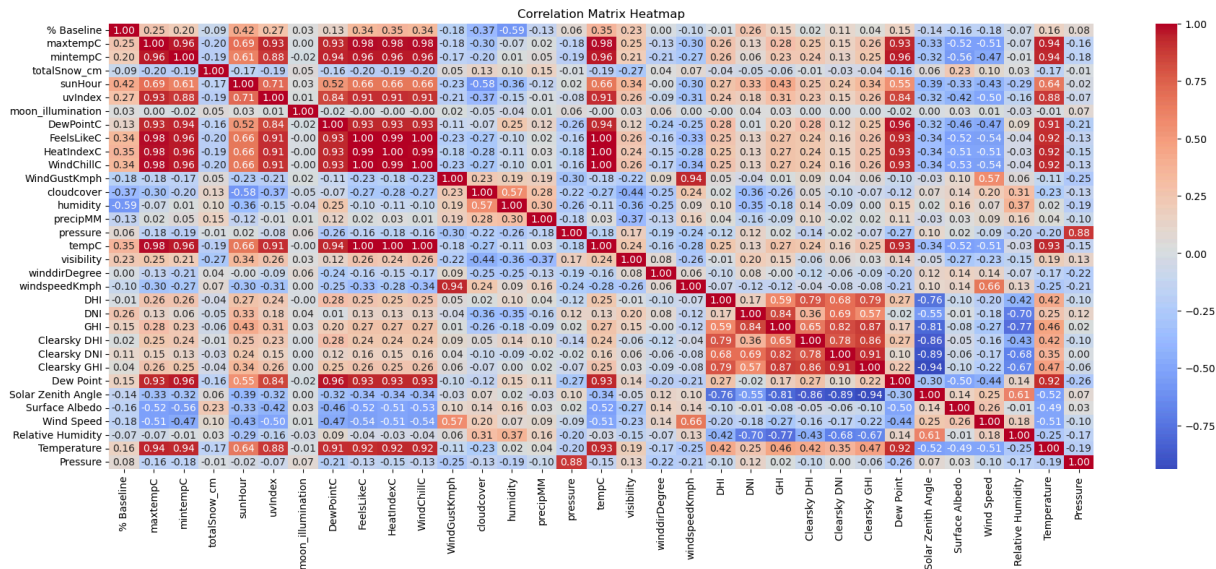
```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\3182376918.py:2: FutureWarning: The default value of `numeric_only` in `DataFrame.corr` is deprecated. In a future version, it will default to `False`. Select only valid columns or specify the value of `numeric_only` to silence this warning.

```

correlation_matrix = join_fil.corr()

```



```
In [8]: baseline_correlation = correlation_matrix['% Baseline']

# Determine which columns to keep (correlation >= 0.2)
columns_to_keep = baseline_correlation[abs(baseline_correlation) >= 0.1].index
exported = join[columns_to_keep]


# Drop columns with correlation < 0.2
filtered_df = join_fil[columns_to_keep]

# Display the filtered DataFrame
filtered_df.head()
```

```
Out[8]:
```

	% Baseline	maxtempC	mintempC	sunHour	uvIndex	DewPointC	FeelsLikeC	HeatIndexC
0	0.0079	-3	-6	8.7	2	-14	-13	-5
1	0.1019	-3	-6	8.7	2	-14	-12	-5
2	0.3932	-3	-6	8.7	2	-14	-11	-4
3	0.5447	-3	-6	8.7	2	-14	-10	-4
4	0.5485	-3	-6	8.7	2	-14	-10	-3

5 rows × 24 columns



```
In [16]: sub = pd.read_csv('dataset/sample_submission.csv')
sub['Timestamp'] = pd.to_datetime(sub['Timestamp'])
sub.head()
```

Out[16]:

	Timestamp	% Baseline
--	-----------	------------

0	2017-10-01 06:00:00	0.0005
1	2017-10-01 07:00:00	0.0005
2	2017-10-01 08:00:00	0.0005
3	2017-10-01 09:00:00	0.0005
4	2017-10-01 10:00:00	0.0005

In [15]: `df_filter.tail()`

Out[15]:

	Timestamp	% Baseline
--	-----------	------------

23263	2017-09-30 14:00:00	0.1846
23264	2017-09-30 15:00:00	0.0711
23265	2017-09-30 16:00:00	0.0560
23266	2017-09-30 17:00:00	0.0182
23267	2017-09-30 18:00:00	0.0004

In [17]: `X.tail()`

Out[17]:

	date_time	maxtempC	mintempC	totalSnow_cm	sunHour	uvIndex	moon_illumina
--	-----------	----------	----------	--------------	---------	---------	---------------

35059	2017-12-31 19:00:00	-9	-13	0.2	8.7	1	
35060	2017-12-31 20:00:00	-9	-13	0.2	8.7	1	
35061	2017-12-31 21:00:00	-9	-13	0.2	8.7	1	
35062	2017-12-31 22:00:00	-9	-13	0.2	8.7	1	
35063	2017-12-31 23:00:00	-9	-13	0.2	8.7	1	

5 rows × 24 columns



In [46]: `df_filter['Timestamp'] = pd.to_datetime(df_filter['Timestamp'])
X['date_time'] = pd.to_datetime(X['date_time'])
full_timestamp_range = pd.date_range(start=df_filter['Timestamp'].min(),`

```

end=X['date_time'].max(),
freq='H') # Assuming hourly frequency
df_extended = pd.DataFrame(full_timestamp_range, columns=['Timestamp'])
df_extended = pd.merge(df_extended, df_filter, on='Timestamp', how='left')
df_extended = pd.merge(df_extended, X, left_on='Timestamp', right_on='date_time', how='left')
df_extended = pd.merge(df_extended, solar, on='Timestamp', how='inner')
df_extended.head()

```

C:\Users\NAUFAL\AppData\Local\Temp\ipykernel_30448\1624733487.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.

Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
df_filter['Timestamp'] = pd.to_datetime(df_filter['Timestamp'])
```

Out[46]:

	Timestamp	% Baseline	date_time	maxtempC	mintempC	totalSnow_cm	sunHour	uvInd
0	2014-01-01 07:00:00	0.0079	2014-01-01 07:00:00	-3	-6	0.0	8.7	
1	2014-01-01 08:00:00	0.1019	2014-01-01 08:00:00	-3	-6	0.0	8.7	
2	2014-01-01 09:00:00	0.3932	2014-01-01 09:00:00	-3	-6	0.0	8.7	
3	2014-01-01 10:00:00	0.5447	2014-01-01 10:00:00	-3	-6	0.0	8.7	
4	2014-01-01 11:00:00	0.5485	2014-01-01 11:00:00	-3	-6	0.0	8.7	

5 rows × 40 columns



```

In [51]: columns_list = list(columns_to_keep) # Convert Index to a List
columns_list.insert(0, 'Timestamp') # Insert 'Timestamp' at the beginning
columns_list = pd.Index(columns_list)

```

In [53]: columns_list

```

Out[53]: Index(['Timestamp', '% Baseline', 'maxtempC', 'mintempC', 'sunHour', 'uvIndex',
'DewPointC', 'FeelsLikeC', 'HeatIndexC', 'WindChillC', 'WindGustKmph',
'cloudcover', 'humidity', 'precipMM', 'tempC', 'visibility',
'windspeedKmph', 'DNI', 'GHI', 'Clearsky DNI', 'Dew Point',
'Solar Zenith Angle', 'Surface Albedo', 'Wind Speed', 'Temperature'],
dtype='object')

```

```
In [54]: aa = df_extended[columns_list]
aa.columns
```

```
Out[54]: Index(['Timestamp', '% Baseline', 'maxtempC', 'mintempC', 'sunHour', 'uvIndex',
               'DewPointC', 'FeelsLikeC', 'HeatIndexC', 'WindChillC', 'WindGustKmph',
               'cloudcover', 'humidity', 'precipMM', 'tempC', 'visibility',
               'windspeedKmph', 'DNI', 'GHI', 'Clearsky DNI', 'Dew Point',
               'Solar Zenith Angle', 'Surface Albedo', 'Wind Speed', 'Temperature'],
              dtype='object')
```

```
In [55]: aa.shape
```

```
Out[55]: (35057, 25)
```





























```
In [56]: aa.to_csv('train_clean_v2.csv')
```

```
In [136... # Assuming `filtered_df` is your DataFrame and '% Baseline' is the target variable
X = filtered_df.drop(columns=['% Baseline'])
y = filtered_df['% Baseline']





























# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_sta

# Standardize the features (important for neural networks)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

```
In [137... ann = Sequential()
unit = (filtered_df.shape[1]-1)*20
ann.add(Dense(units = unit, activation = "relu"))
ann.add(Dense(units = unit, activation = "relu"))
ann.add(Dense(units = unit, activation = "relu"))
ann.add(Dense(units = unit, activation = "relu"))
ann.add(Dense(units = 1))
ann.compile(optimizer = "adam", loss = "mean_squared_error")
early_stop = EarlyStopping(monitor='val_loss', mode='min', patience=20, restore_best
history = ann.fit(X_train, y_train, batch_size = 32, epochs = 150, validation_data=
```


Epoch 1/150
421/421  15s 24ms/step - loss: 0.0429 - val_loss: 0.0170
Epoch 2/150
421/421  10s 25ms/step - loss: 0.0161 - val_loss: 0.0171
Epoch 3/150
421/421  20s 24ms/step - loss: 0.0140 - val_loss: 0.0137
Epoch 4/150
421/421  11s 27ms/step - loss: 0.0131 - val_loss: 0.0136
Epoch 5/150
421/421  12s 29ms/step - loss: 0.0120 - val_loss: 0.0134
Epoch 6/150
421/421  20s 27ms/step - loss: 0.0116 - val_loss: 0.0120
Epoch 7/150
421/421  10s 24ms/step - loss: 0.0107 - val_loss: 0.0122
Epoch 8/150
421/421  9s 21ms/step - loss: 0.0102 - val_loss: 0.0115
Epoch 9/150
421/421  10s 23ms/step - loss: 0.0094 - val_loss: 0.0132
Epoch 10/150
421/421  11s 24ms/step - loss: 0.0094 - val_loss: 0.0111
Epoch 11/150
421/421  11s 25ms/step - loss: 0.0087 - val_loss: 0.0127
Epoch 12/150
421/421  10s 24ms/step - loss: 0.0087 - val_loss: 0.0107
Epoch 13/150
421/421  8s 19ms/step - loss: 0.0072 - val_loss: 0.0095
Epoch 14/150
421/421  8s 19ms/step - loss: 0.0072 - val_loss: 0.0112
Epoch 15/150
421/421  10s 23ms/step - loss: 0.0069 - val_loss: 0.0098
Epoch 16/150
421/421  10s 23ms/step - loss: 0.0061 - val_loss: 0.0099
Epoch 17/150
421/421  10s 23ms/step - loss: 0.0059 - val_loss: 0.0110
Epoch 18/150
421/421  10s 23ms/step - loss: 0.0063 - val_loss: 0.0088
Epoch 19/150
421/421  10s 22ms/step - loss: 0.0058 - val_loss: 0.0092
Epoch 20/150
421/421  9s 22ms/step - loss: 0.0054 - val_loss: 0.0098
Epoch 21/150
421/421  9s 22ms/step - loss: 0.0048 - val_loss: 0.0096
Epoch 22/150
421/421  11s 24ms/step - loss: 0.0046 - val_loss: 0.0089
Epoch 23/150
421/421  11s 25ms/step - loss: 0.0044 - val_loss: 0.0087
Epoch 24/150
421/421  9s 22ms/step - loss: 0.0041 - val_loss: 0.0088
Epoch 25/150
421/421  11s 26ms/step - loss: 0.0040 - val_loss: 0.0084
Epoch 26/150
421/421  11s 26ms/step - loss: 0.0036 - val_loss: 0.0085
Epoch 27/150
421/421  10s 23ms/step - loss: 0.0040 - val_loss: 0.0084
Epoch 28/150
421/421  9s 22ms/step - loss: 0.0035 - val_loss: 0.0081

Epoch 29/150
421/421 ————— 9s 22ms/step - loss: 0.0035 - val_loss: 0.0080
Epoch 30/150
421/421 ————— 9s 22ms/step - loss: 0.0033 - val_loss: 0.0083
Epoch 31/150
421/421 ————— 11s 23ms/step - loss: 0.0030 - val_loss: 0.0079
Epoch 32/150
421/421 ————— 9s 22ms/step - loss: 0.0028 - val_loss: 0.0078
Epoch 33/150
421/421 ————— 11s 24ms/step - loss: 0.0025 - val_loss: 0.0085
Epoch 34/150
421/421 ————— 10s 23ms/step - loss: 0.0029 - val_loss: 0.0080
Epoch 35/150
421/421 ————— 9s 22ms/step - loss: 0.0028 - val_loss: 0.0085
Epoch 36/150
421/421 ————— 9s 22ms/step - loss: 0.0026 - val_loss: 0.0074
Epoch 37/150
421/421 ————— 11s 23ms/step - loss: 0.0027 - val_loss: 0.0076
Epoch 38/150
421/421 ————— 10s 23ms/step - loss: 0.0027 - val_loss: 0.0076
Epoch 39/150
421/421 ————— 10s 24ms/step - loss: 0.0021 - val_loss: 0.0076
Epoch 40/150
421/421 ————— 10s 23ms/step - loss: 0.0022 - val_loss: 0.0081
Epoch 41/150
421/421 ————— 9s 22ms/step - loss: 0.0020 - val_loss: 0.0083
Epoch 42/150
421/421 ————— 9s 22ms/step - loss: 0.0021 - val_loss: 0.0081
Epoch 43/150
421/421 ————— 10s 22ms/step - loss: 0.0024 - val_loss: 0.0074
Epoch 44/150
421/421 ————— 9s 22ms/step - loss: 0.0019 - val_loss: 0.0074
Epoch 45/150
421/421 ————— 9s 22ms/step - loss: 0.0017 - val_loss: 0.0074
Epoch 46/150
421/421 ————— 9s 22ms/step - loss: 0.0016 - val_loss: 0.0074
Epoch 47/150
421/421 ————— 9s 22ms/step - loss: 0.0018 - val_loss: 0.0074
Epoch 48/150
421/421 ————— 9s 22ms/step - loss: 0.0020 - val_loss: 0.0071
Epoch 49/150
421/421 ————— 10s 23ms/step - loss: 0.0016 - val_loss: 0.0072
Epoch 50/150
421/421 ————— 9s 22ms/step - loss: 0.0018 - val_loss: 0.0075
Epoch 51/150
421/421 ————— 9s 22ms/step - loss: 0.0018 - val_loss: 0.0072
Epoch 52/150
421/421 ————— 10s 22ms/step - loss: 0.0016 - val_loss: 0.0072
Epoch 53/150
421/421 ————— 9s 21ms/step - loss: 0.0018 - val_loss: 0.0071
Epoch 54/150
421/421 ————— 9s 22ms/step - loss: 0.0016 - val_loss: 0.0071
Epoch 55/150
421/421 ————— 10s 22ms/step - loss: 0.0015 - val_loss: 0.0069
Epoch 56/150
421/421 ————— 9s 22ms/step - loss: 0.0012 - val_loss: 0.0068

Epoch 57/150
421/421  10s 22ms/step - loss: 0.0013 - val_loss: 0.0070
Epoch 58/150
421/421  9s 22ms/step - loss: 0.0012 - val_loss: 0.0070
Epoch 59/150
421/421  9s 22ms/step - loss: 0.0015 - val_loss: 0.0073
Epoch 60/150
421/421  9s 22ms/step - loss: 0.0013 - val_loss: 0.0069
Epoch 61/150
421/421  9s 20ms/step - loss: 0.0012 - val_loss: 0.0069
Epoch 62/150
421/421  8s 18ms/step - loss: 0.0012 - val_loss: 0.0074
Epoch 63/150
421/421  8s 20ms/step - loss: 0.0012 - val_loss: 0.0072
Epoch 64/150
421/421  8s 20ms/step - loss: 0.0012 - val_loss: 0.0070
Epoch 65/150
421/421  8s 20ms/step - loss: 0.0012 - val_loss: 0.0078
Epoch 66/150
421/421  8s 18ms/step - loss: 0.0015 - val_loss: 0.0071
Epoch 67/150
421/421  11s 19ms/step - loss: 0.0011 - val_loss: 0.0072
Epoch 68/150
421/421  8s 20ms/step - loss: 0.0010 - val_loss: 0.0070
Epoch 69/150
421/421  10s 20ms/step - loss: 0.0011 - val_loss: 0.0067
Epoch 70/150
421/421  8s 20ms/step - loss: 0.0010 - val_loss: 0.0067
Epoch 71/150
421/421  8s 20ms/step - loss: 0.0012 - val_loss: 0.0076
Epoch 72/150
421/421  11s 26ms/step - loss: 0.0017 - val_loss: 0.0069
Epoch 73/150
421/421  9s 21ms/step - loss: 9.2031e-04 - val_loss: 0.0067
Epoch 74/150
421/421  9s 21ms/step - loss: 7.0434e-04 - val_loss: 0.0069
Epoch 75/150
421/421  9s 21ms/step - loss: 8.3645e-04 - val_loss: 0.0073
Epoch 76/150
421/421  9s 21ms/step - loss: 0.0012 - val_loss: 0.0069
Epoch 77/150
421/421  9s 22ms/step - loss: 9.4853e-04 - val_loss: 0.0069
Epoch 78/150
421/421  10s 22ms/step - loss: 9.2484e-04 - val_loss: 0.0068
Epoch 79/150
421/421  10s 21ms/step - loss: 8.2835e-04 - val_loss: 0.0081
Epoch 80/150
421/421  10s 21ms/step - loss: 0.0015 - val_loss: 0.0068
Epoch 81/150
421/421  9s 22ms/step - loss: 8.3470e-04 - val_loss: 0.0069
Epoch 82/150
421/421  9s 21ms/step - loss: 8.7553e-04 - val_loss: 0.0068
Epoch 83/150
421/421  10s 21ms/step - loss: 8.8971e-04 - val_loss: 0.0069
Epoch 84/150
421/421  11s 21ms/step - loss: 7.7441e-04 - val_loss: 0.0069

```

Epoch 85/150
421/421 ————— 9s 22ms/step - loss: 8.1930e-04 - val_loss: 0.0068
Epoch 86/150
421/421 ————— 9s 21ms/step - loss: 7.8230e-04 - val_loss: 0.0065
Epoch 87/150
421/421 ————— 10s 21ms/step - loss: 6.5902e-04 - val_loss: 0.0069
Epoch 88/150
421/421 ————— 10s 21ms/step - loss: 7.8525e-04 - val_loss: 0.0071
Epoch 89/150
421/421 ————— 9s 22ms/step - loss: 0.0011 - val_loss: 0.0068
Epoch 90/150
421/421 ————— 10s 22ms/step - loss: 7.9105e-04 - val_loss: 0.0068
Epoch 91/150
421/421 ————— 10s 23ms/step - loss: 6.4908e-04 - val_loss: 0.0072
Epoch 92/150
421/421 ————— 10s 22ms/step - loss: 0.0011 - val_loss: 0.0073
Epoch 93/150
421/421 ————— 9s 22ms/step - loss: 8.8328e-04 - val_loss: 0.0068
Epoch 94/150
421/421 ————— 10s 21ms/step - loss: 6.7517e-04 - val_loss: 0.0066
Epoch 95/150
421/421 ————— 9s 21ms/step - loss: 8.8136e-04 - val_loss: 0.0079
Epoch 96/150
421/421 ————— 10s 21ms/step - loss: 0.0015 - val_loss: 0.0070
Epoch 97/150
421/421 ————— 9s 22ms/step - loss: 6.6023e-04 - val_loss: 0.0067
Epoch 98/150
421/421 ————— 9s 22ms/step - loss: 4.5939e-04 - val_loss: 0.0068
Epoch 99/150
421/421 ————— 9s 21ms/step - loss: 4.5979e-04 - val_loss: 0.0067
Epoch 100/150
421/421 ————— 9s 22ms/step - loss: 4.4530e-04 - val_loss: 0.0069
Epoch 101/150
421/421 ————— 10s 21ms/step - loss: 5.2688e-04 - val_loss: 0.0068
Epoch 102/150
421/421 ————— 11s 22ms/step - loss: 6.8092e-04 - val_loss: 0.0072
Epoch 103/150
421/421 ————— 9s 22ms/step - loss: 9.6327e-04 - val_loss: 0.0078
Epoch 104/150
421/421 ————— 9s 22ms/step - loss: 0.0015 - val_loss: 0.0068
Epoch 105/150
421/421 ————— 9s 22ms/step - loss: 6.9048e-04 - val_loss: 0.0067
Epoch 106/150
421/421 ————— 9s 22ms/step - loss: 5.1529e-04 - val_loss: 0.0071

```


In [138...

```

# Plot training & validation loss values
plt.figure(figsize=(6, 5))
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Test'], loc='upper left')

plt.tight_layout()
plt.show()

```



```
In [139... pred = ann.predict(X_test)
tpred = ann.predict(X_train)
print('Train R2 Score : ',r2_score(y_train,tpred))
print('Test R2 Score : ',r2_score(y_test,pred))
```

```
106/106 ————— 1s 9ms/step
421/421 ————— 1s 3ms/step
Train R2 Score : 0.990520082179919
Test R2 Score : 0.9017448124976732
```

```
In [140... 0.90-0.895
```


```
Out[140... 0.00500000000000000044
```

```
In [134... # Save the entire model
#ann.save('ann_model.h5')
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```
In [56]: plt.figure(figsize=(15,6))
plt.grid(True)
plt.xlabel('Time')
plt.ylabel('%Baseline')
```


```
plt.plot(df_filter['% Baseline'])
plt.title('NPHR Time Series')
plt.show()
```



```
In [57]: #Test for stationarity
def test_stationarity(timeseries):
    #Determining rolling statistics
    rolmean = timeseries.rolling(12).mean()
    rolstd = timeseries.rolling(12).std()
    #Plot rolling statistics:
    plt.plot(timeseries, color='blue',label='Original')
    plt.plot(rolmean, color='red', label='Rolling Mean')
    plt.plot(rolstd, color='black', label = 'Rolling Std')
    plt.legend(loc='best')
    plt.title('Rolling Mean and Standard Deviation')
    plt.show(block=False)

    print("Results of Dickey Fuller test")
    adft = adfuller(timeseries,autolag='AIC')
    # output for dft will give us without defining what the values are.
    #hence we manually write what values does it explains using a for Loop
    output = pd.Series(adft[0:4],index=['Test Statistics','p-value','No. of lags us
    for key,values in adft[4].items():
        output['critical value (%s)'%key] = values
    print(output)

test_stationarity(df_filter['% Baseline'])
```

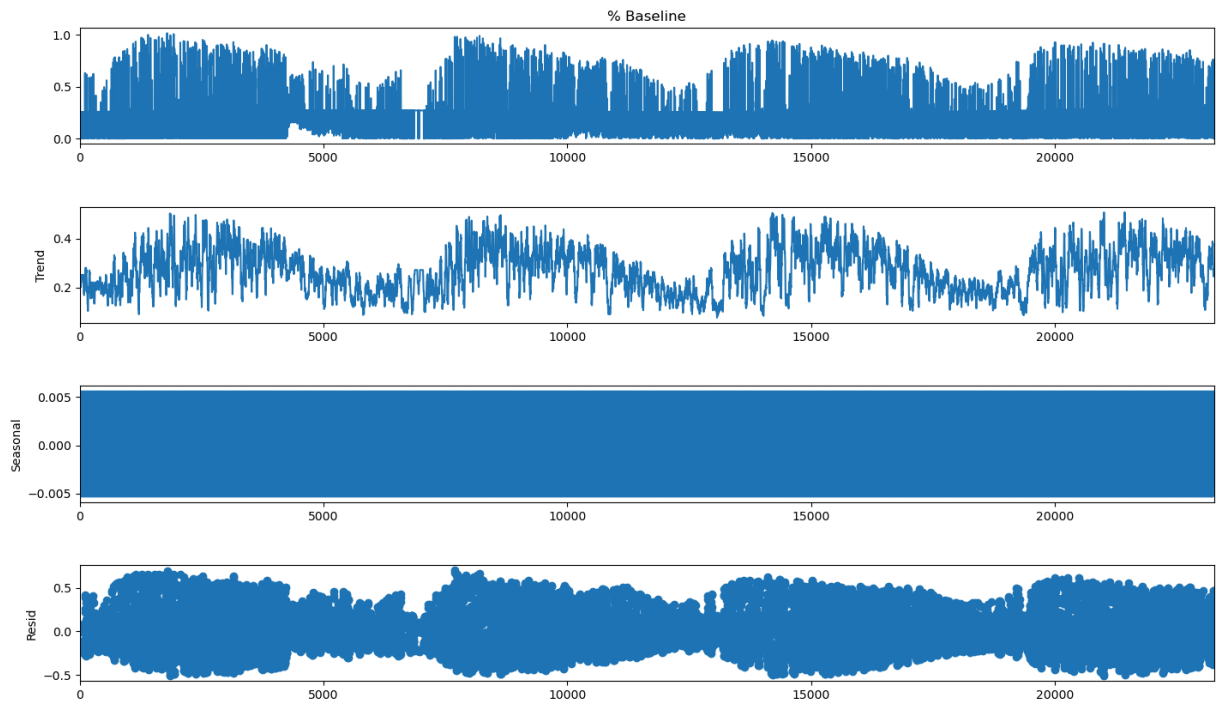


Results of Dickey Fuller test

Test Statistics	-1.269716e+01
p-value	1.099764e-23
No. of lags used	4.700000e+01
Number of observations used	2.322000e+04
critical value (1%)	-3.430632e+00
critical value (5%)	-2.861664e+00
critical value (10%)	-2.566836e+00
dtype:	float64

```
In [58]: result = seasonal_decompose(df_filter['% Baseline'], model='additive', period=30)
fig = plt.figure()
fig = result.plot()
fig.set_size_inches(16, 9)
```

<Figure size 640x480 with 0 Axes>



```
In [18]: XX = X_filtered[['WindChillC', 'WindGustKmph', 'cloudcover']]
model_autoARIMA = auto_arima(df_filter['% Baseline'], start_p=0, start_q=0,
                             test='adf',          # use adftest to find optimal
                             max_p=3, max_q=3,    # maximum p and q
                             m=1,                # frequency of series
                             d=None,              # Let model determine 'd'
                             seasonal=False,      # No Seasonality
                             start_P=0,
                             D=0,
                             trace=True,
                             error_action='ignore',
                             suppress_warnings=True,
                             stepwise=True,
                             exogenous= XX)
print(model_autoARIMA.summary())
```


Performing stepwise search to minimize aic

```
ARIMA(0,0,0)(0,0,0)[0] : AIC=12924.902, Time=1.09 sec
ARIMA(1,0,0)(0,0,0)[0] : AIC=-40558.901, Time=0.75 sec
ARIMA(0,0,1)(0,0,0)[0] : AIC=-13745.822, Time=2.16 sec
ARIMA(2,0,0)(0,0,0)[0] : AIC=-51223.708, Time=1.48 sec
ARIMA(3,0,0)(0,0,0)[0] : AIC=-51230.933, Time=2.34 sec
ARIMA(3,0,1)(0,0,0)[0] : AIC=-51230.235, Time=10.96 sec
ARIMA(2,0,1)(0,0,0)[0] : AIC=-51230.333, Time=3.35 sec
ARIMA(3,0,0)(0,0,0)[0] intercept : AIC=-53351.572, Time=5.90 sec
ARIMA(2,0,0)(0,0,0)[0] intercept : AIC=-53093.955, Time=1.75 sec
ARIMA(3,0,1)(0,0,0)[0] intercept : AIC=-53414.515, Time=29.78 sec
ARIMA(2,0,1)(0,0,0)[0] intercept : AIC=-53424.572, Time=13.09 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=-48823.340, Time=6.41 sec
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=-53515.518, Time=24.85 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=-50926.113, Time=16.13 sec
ARIMA(3,0,2)(0,0,0)[0] intercept : AIC=-53460.738, Time=48.57 sec
ARIMA(2,0,3)(0,0,0)[0] intercept : AIC=-53508.371, Time=30.87 sec
ARIMA(1,0,3)(0,0,0)[0] intercept : AIC=-51914.036, Time=18.62 sec
ARIMA(3,0,3)(0,0,0)[0] intercept : AIC=-53511.728, Time=30.65 sec
ARIMA(2,0,2)(0,0,0)[0] : AIC=-51237.738, Time=5.49 sec
```

Best model: ARIMA(2,0,2)(0,0,0)[0] intercept

Total fit time: 254.365 seconds

SARIMAX Results

```
=====
Dep. Variable:          y      No. Observations:      23268
Model:                 SARIMAX(2, 0, 2)  Log Likelihood      26763.759
Date:                 Tue, 20 Aug 2024    AIC              -53515.518
Time:                 19:00:53           BIC              -53467.189
Sample:              0                HQIC             -53499.821
                        - 23268
```

Covariance Type: opg

```
=====
              coef      std err          z      P>|z|      [0.025      0.975]
-----
intercept      0.0249      0.001     31.475      0.000      0.023      0.026
ar.L1          1.6595      0.008    199.448      0.000      1.643      1.676
ar.L2         -0.7806      0.007   -109.364      0.000     -0.795     -0.767
ma.L1         -0.2444      0.009    -25.802      0.000     -0.263     -0.226
ma.L2         -0.0801      0.007    -11.118      0.000     -0.094     -0.066
sigma2         0.0059    3.57e-05    164.100      0.000      0.006      0.006
=====
```

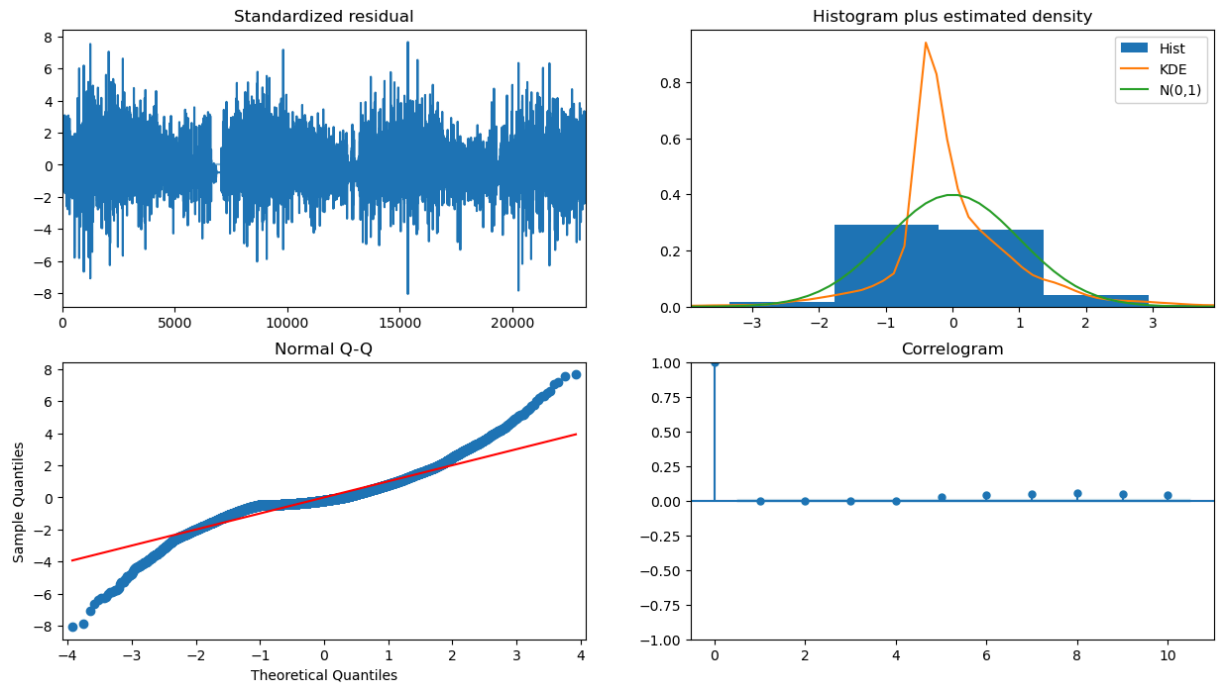
```
=====
Ljung-Box (L1) (Q):      0.00  Jarque-Bera (JB):      32858.16
Prob(Q):                0.99  Prob(JB):                0.00
Heteroskedasticity (H):  0.95  Skew:              0.50
Prob(H) (two-sided):    0.02  Kurtosis:          8.73
=====
```

Warnings:

[1] Covariance matrix calculated using the outer product of gradients (complex-step).


```
In [19]: def mape(actual, pred):
         actual, pred = np.array(actual), np.array(pred)
         return np.mean(np.abs((actual - pred) / actual)) * 100
```

```
In [20]: model_autoARIMA.plot_diagnostics(figsize=(15,8))
plt.show()
```




```
In [32]: fit_bs = model_autoARIMA.fittedvalues()
plt.figure(figsize=(10,6))
plt.grid(True)
plt.xlabel('Dates')
plt.ylabel('NPHR')
plt.plot(df_filter['% Baseline'], 'green', label='Actual Baseline')
plt.plot(fit_bs, 'red', label='Fit Baseline')
plt.legend()
print(np.sqrt(mean_squared_error(df_filter['% Baseline'], fit_bs)))
```

0.07660454589992081



```
In [33]: # Replace negative values in fit_bs with 0
fit_bs = np.where(fit_bs < 0, 0, fit_bs)

# Plot the results again
plt.figure(figsize=(10,6))
plt.grid(True)
plt.xlabel('Dates')
plt.ylabel('NPHR')
plt.plot(df_filter['% Baseline'], 'green', label='Actual Baseline')
plt.plot(fit_bs, 'red', label='Fit Baseline', alpha=0.2)
plt.legend()
plt.show()
print(np.sqrt(mean_squared_error(df_filter['% Baseline'], fit_bs)))
```



0.07564483165475702